




Synthesis of Sorting Kernels

03.03.2025, CGO 2025


Marcel Ullrich, Sebastian Hack

Saarland University, Saarland Informatics Campus


Sorting Kernels



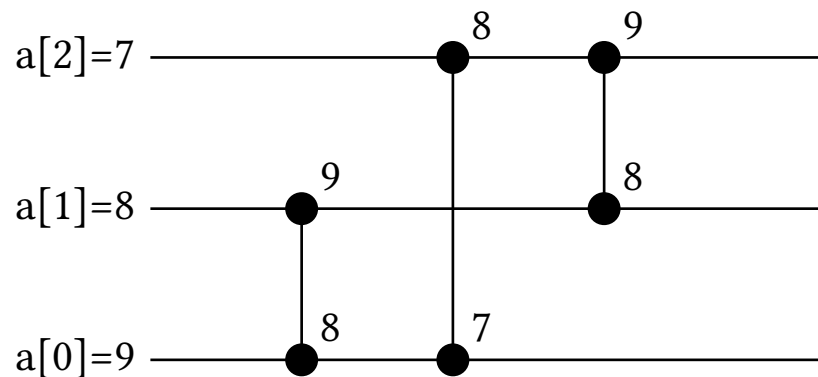

Sorting Kernels



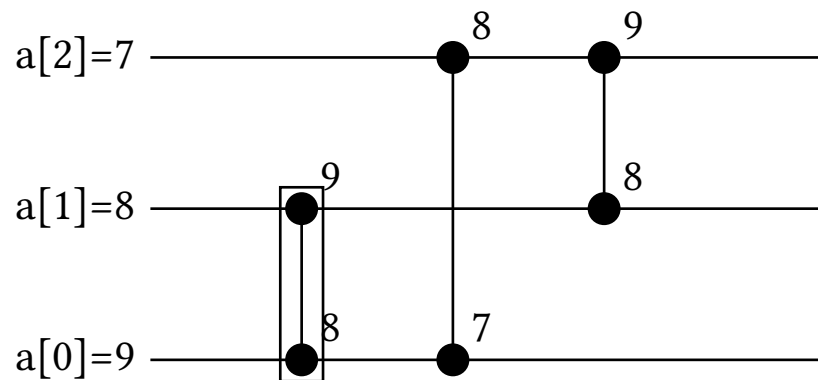

Sorting Kernels



Sorting Kernels



Sorting Kernels



```
mov rdi, rax  
cmp rbx, rax  
cmovl rax, rbx  
cmovl rbx, rdi
```



Model

| | register | | swap | lt | gt |
|---|----------|----|------|----|----|
| mov s1 r2 | 13 | 9 | - | - | - |
| cmp r1 r2 | 13 | 9 | 9 | - | - |
| cmovg r2 r1 | 13 | 9 | 9 | - | > |
| cmovg r1 s1 | 13 | 13 | 9 | - | > |
| sample from {mov, cmovl, cmovg, cmp} | 9 | 13 | 9 | - | > |



Search Space

| | register | | swap | lt | gt |
|---|----------|----|------|----|----|
| mov s1 r2 | 13 | 9 | - | - | - |
| cmp r1 r2 | 13 | 9 | 9 | - | - |
| cmovg r2 r1 | 13 | 9 | 9 | - | > |
| cmovg r1 s1 | 13 | 13 | 9 | - | > |
| sample from {mov, cmovl, cmovg, cmp} | 9 | 13 | 9 | - | > |

$\min(a, \min(b, c)) =$
 $\min(\min(\max(c, b), a), \min(b, c))$



Search Space

mov s1 r2

cmp r1 r2

cmovg r2 r1

cmovg r1 s1

sample from
{mov, cmovl, cmovg, cmp}

| register | swap | lt | gt |
|----------|------|----|----|
| 13 9 | - | - | - |
| 13 9 | 9 | - | - |
| 13 9 | 9 | - | > |
| 13 13 | 9 | - | > |
| 9 13 | 9 | - | > |

| n | Program Length | Search Space |
|-----|----------------|--------------|
| 3 | 11 | $10^{19.9}$ |
| 4 | 20 | 10^{40} |
| 5 | ≈ 33 | $10^{71.2}$ |
| 6 | ≈ 45 | $10^{108.4}$ |

$$\min(a, \min(b, c)) = \min(\min(\max(c, b), a), \min(b, c))$$

5602 solutions for $n = 3$



Search Space

| | register | swap | lt | gt |
|-------------|----------|------|----|----|
| mov s1 r2 | 13 9 | - | - | - |
| cmp r1 r2 | 13 9 | 9 | - | - |
| cmovg r2 r1 | 13 9 | 9 | - | > |
| cmovg r1 s1 | 13 13 | 9 | - | > |
| | 9 13 | 9 | - | > |

sample from
{mov, cmovl, cmovg, cmp}



TSNE-embedding of solutions

| n | Program Length | Search Space |
|-----|----------------|--------------|
| 3 | 11 | $10^{19.9}$ |
| 4 | 20 | 10^{40} |
| 5 | ≈ 33 | $10^{71.2}$ |
| 6 | ≈ 45 | $10^{108.4}$ |

5602 solutions for $n = 3$



State of the Art

- sorting network 🐌
-
-





State of the Art

- sorting network 🐌
- handoptimized 🐛
-





State of the Art

- sorting network 
- handoptimized 
- 2024 AlphaDev¹
 - $n = 3$: 6min
 - $n = 4$: 30min
 - $n = 5$: 17.5h

¹ Mankowitz, Daniel J., et al. “Faster sorting algorithms discovered using deep reinforcement learning.” Nature 618. (2023): 257-263.





State of the Art

- sorting network 
- handoptimized 
- 2024 AlphaDev
 - $n = 3$: ~~6min~~ 97ms
 - $n = 4$: ~~30min~~ 2.4s
 - $n = 5$: ~~17.5h~~ 11min

 faster synthesis



State of the Art

- sorting network 
- handoptimized 
- 2024 AlphaDev
 - $n = 3$: ~~6min~~ 97ms
 - $n = 4$: ~~30min~~ 2.4s
 - $n = 5$: ~~17.5h~~ 11min



faster synthesis



faster sorting kernels




State of the Art

- sorting network 🐌
- handoptimized 🐛
- 2024 AlphaDev
 - ▶ $n = 3$: ~~6min~~ 97ms
 - ▶ $n = 4$: ~~30min~~ 2.4s
 - ▶ $n = 5$: ~~17.5h~~ 11min


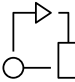




- 🏃 faster synthesis
- 📈 faster sorting kernels
- 📦 minimality proof



Enumerative Synthesis



Enumerative Synthesis


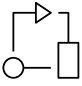




1.  Select open state
2.  Apply Instruction
3.  Check for viability
4.  Check for solution
5.  Cut non-promising
6.  Deduplicate states

A★ with heuristics:

- permutations
- permutations + scratch register
- delete-relaxed
(maximum per permutation)



Enumerative Synthesis

1.  Select open state
2.  Apply Instruction
3.  Check for viability
4.  Check for solution
5.  Cut non-promising
6.  Deduplicate states

Remove redundant/non-sensical:


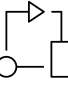




- `cmp r1 r2; cmp r1 r3`
- `cmp r1 r1`

Restrict to beneficial:

- `delete-relaxed`
- `cmp r2 r1 → cmp r1 r2`



Enumerative Synthesis


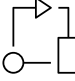




1.  Select open state
2.  Apply Instruction
3.  Check for viability
4.  Check for solution
5.  Cut non-promising
6.  Deduplicate states

Cut programs:

- number eliminated
- longer than bound/solution
- can not be completed in time




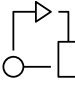




Enumerative Synthesis

1.  Select open state
2.  Apply Instruction
3.  Check for viability
4.  Check for solution
5.  Cut non-promising
6.  Deduplicate states

All permutations already sorted



Enumerative Synthesis

1.  Select open state
2.  Apply Instruction
3.  Check for viability
4.  Check for solution
5.  Cut non-promising
6.  Deduplicate states


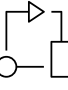




Cut if

permutation count $> k \times \text{best}$

| Cut | Solutions |
|--------------|-----------|
| $k = 1$ | 222 |
| $k = 1.5$ | 838 |
| $k = 2$ | 5602 |
| $k = \infty$ | 5602 |



Enumerative Synthesis

1.  Select open state
2.  Apply Instruction
3.  Check for viability
4.  Check for solution
5.  Cut non-promising
6.  Deduplicate states

Hashset-based deduplication
of states



Solver-Based Techniques

$$\begin{aligned} & \forall r : P(r) = o \rightarrow \\ & \underbrace{(\forall 1 \leq i \leq |r| : o_i \leq o_{i+1})}_{\text{ascending}} \wedge \\ & \underbrace{(\forall x : |\{i : r_i = x\}| = |\{i : o_i = x\}|)}_{\text{same elements}} \end{aligned}$$



Solver-Based Techniques

$$r \in \text{Perm}(1..n)$$

$$\forall r : P(r) = o \rightarrow \forall 1 \leq i \leq r : o_i = i$$



Solver-Based Techniques

$$\bigwedge_{r \in \text{Perm}(1..n)} \bigwedge_{1 \leq i \leq n} P(r)_i = i$$



Solver-Based Techniques

$$\bigwedge_{r \in \text{Perm}(1..n)} \bigwedge_{1 \leq i \leq n} P(r)_i = i$$

Heuristics:

- `cmp r1 r2; cmp r2 r3` \rightarrow `cmp r2 r3`
- `cmp r1 r1` \rightarrow `noop`
- `cmp r3 r2` \rightarrow `cmp r2 r3`
- only read initialized
- do not make uncompleteable



Solver-Based Synthesis $n = 3$

| SMT | Approach |
|-------|-------------------------|
| 97min | CEGIS, arbitrary inputs |
| 25min | CEGIS, 1..n |
| 44min | all permutations |
| — | SyGuS (CVC5, Metalift) |



Solver-Based Synthesis $n = 3$

| SMT | Approach | CP | Approach |
|-------|-------------------------|------|--------------------------|
| 97min | CEGIS, arbitrary inputs | — | ILP, MIP |
| 25min | CEGIS, 1.. n | — | CP (MiniZinc other) |
| 44min | all permutations | 232s | chuffed, no heuristic |
| — | SyGuS (CVC5, Metalift) | 70s | chuffed, $h, = 1..n$ |
| | | 30s | chuffed, $h, \leq, 1..3$ |



Solver-Based Synthesis $n = 3$

| SMT | Approach |
|-------|-------------------------|
| 97min | CEGIS, arbitrary inputs |
| 25min | CEGIS, 1.. n |
| 44min | all permutations |
| — | SyGuS (CVC5, Metalift) |

| Planning | Approach |
|----------|-----------------------|
| 679s | Scorpion planner |
| 216s | Lama planner grounded |
| 3.54s | Lama Planner |

| CP | Approach |
|------|--------------------------|
| — | ILP, MIP |
| — | CP (MiniZinc other) |
| 232s | chuffed, no heuristic |
| 70s | chuffed, $h, = 1..n$ |
| 30s | chuffed, $h, \leq, 1..3$ |



Solver-Based Synthesis $n = 3$

| SMT | Approach |
|-------|-------------------------|
| 97min | CEGIS, arbitrary inputs |
| 25min | CEGIS, 1.. n |
| 44min | all permutations |
| — | SyGuS (CVC5, Metalift) |

| Planning | Approach |
|----------|-----------------------|
| 679s | Scorpion planner |
| 216s | Lama planner grounded |
| 3.54s | Lama Planner |

| CP | Approach |
|------|--------------------------|
| — | ILP, MIP |
| — | CP (MiniZinc other) |
| 232s | chuffed, no heuristic |
| 70s | chuffed, $h, = 1..n$ |
| 30s | chuffed, $h, \leq, 1..3$ |

Enum: 97ms



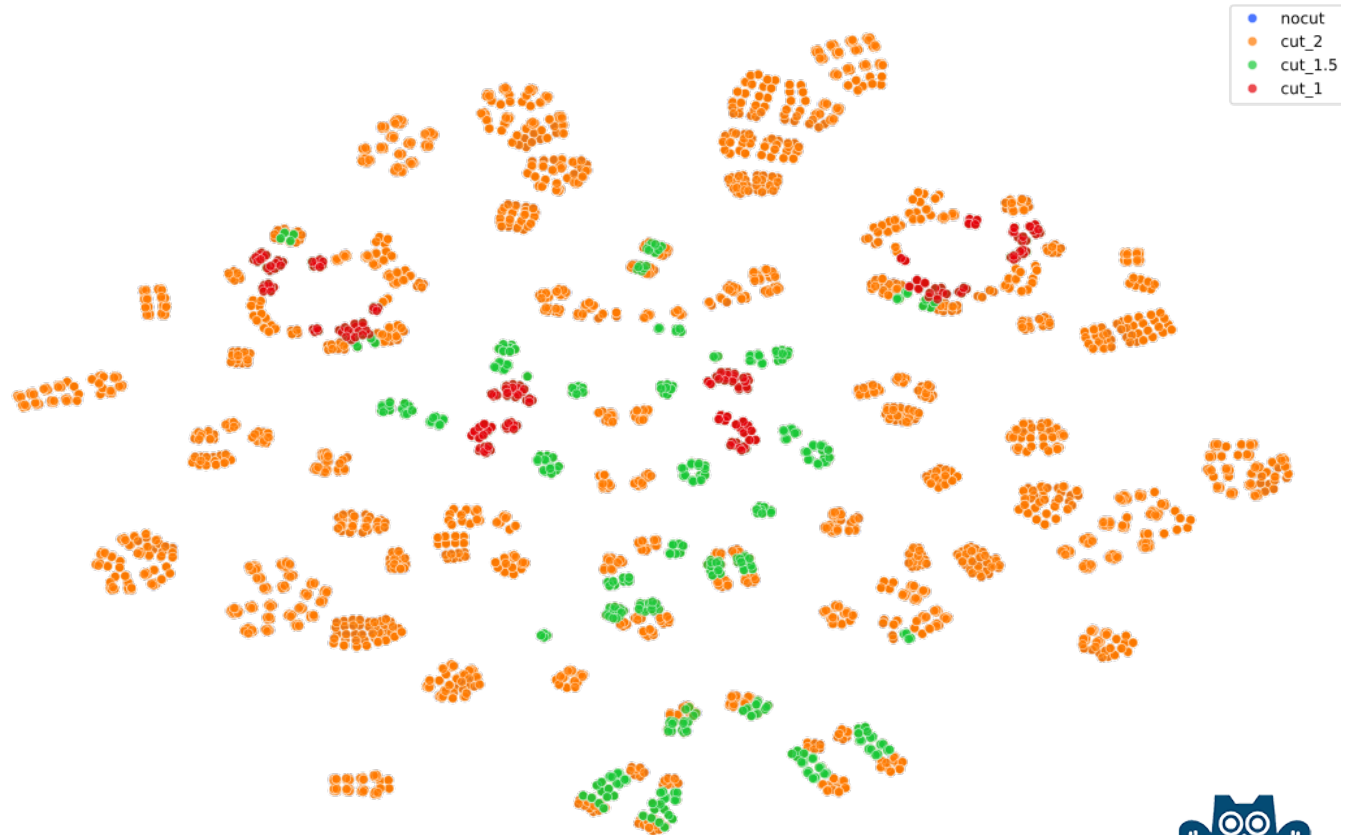
Evaluation Enumeration $n = 3$

| Approach | Time |
|-------------------|------|
| Dijkstra | 56s |
| Dijkstra parallel | 17s |
| Dedup, viable | 8.6s |
| Dedup, A★ | 1.7s |
| +viable, instr | 0.7s |
| +cut $k = 1$ | 0.1s |



Evaluation Enumeration $n = 3$

| Approach | Time |
|-------------------|------|
| Dijkstra | 56s |
| Dijkstra parallel | 17s |
| Dedup, viable | 8.6s |
| Dedup, A★ | 1.7s |
| +viable, instr | 0.7s |
| +cut $k = 1$ | 0.1s |



Evaluation Enumeration $n \geq 3$

| | $l = 11$ | $l = 20$ | $l = 33$ |
|-------------|----------|----------|----------|
| Approach | $n = 3$ | $n = 4$ | $n = 5$ |
| Enumeration | 97ms | 2.4s | 11min |
| AlphaDev-RL | 6min | 30min | 17.5h |
| AlphaDev-S | 0.4s | 0.6s | 5.75h |

- All solutions for $n = 3$: 10min
- Optimality for $n = 4$: 2weeks



Evaluation Kernels

| Kernel | $n = 3$ | $n = 4$ | $n = 5$ |
|-----------------------|---------|---------|---------|
| <u>Enumeration</u> | 5.8ms | 9.4ms | 14.8ms |
| Mimicry ² | 8.0ms | 8.8ms | — |
| AlphaDev | 6.7ms | 10.4ms | 16.2ms |
| Sorting Network (Cmp) | 7.1ms | 14.8ms | 19.4ms |

²Mimicry. 2023. Faster Sorting Beyond DeepMind's AlphaDev. <https://www.mimicry.ai/faster-sorting-beyond-deepminds-alphadev> Accessed: 2023-09-20



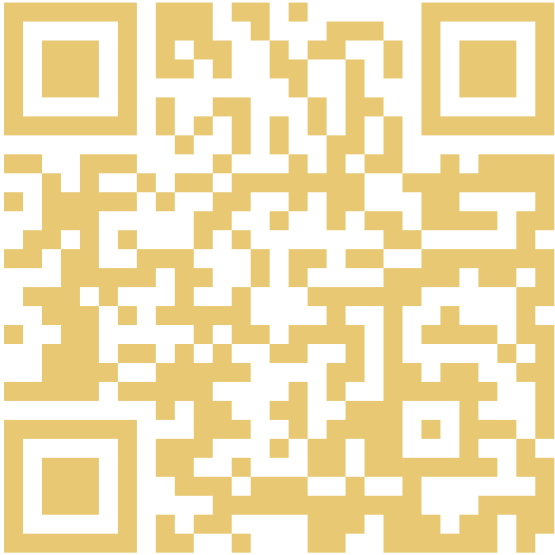
Evaluation Kernels MinMax

| Kernel | $n = 3$ | $n = 4$ | $n = 5$ | |
|-----------------------|---------|---------|---------|---------------------|
| <u>Enumeration</u> | 5.8ms | 9.4ms | 14.8ms | |
| Mimicry ³ | 8.0ms | 8.8ms | — | movdqa %xmm1, %xmm3 |
| AlphaDev | 6.7ms | 10.4ms | 16.2ms | pminud %xmm2, %xmm1 |
| Sorting Network (Cmp) | 7.1ms | 14.8ms | 19.4ms | pmaxud %xmm3, %xmm2 |
| <u>MinMax</u> | 4.6ms | 7.0ms | 10.7ms | movdqa %xmm0, %xmm3 |
| Sorting Network | 5.3ms | 8.1ms | 12.2ms | pminud %xmm2, %xmm3 |
| | | | | pmaxud %xmm0, %xmm2 |
| | | | | pminud %xmm1, %xmm0 |
| | | | | pmaxud %xmm3, %xmm1 |

³Mimicry. 2023. Faster Sorting Beyond DeepMind's AlphaDev. <https://www.mimicry.ai/faster-sorting-beyond-deepminds-alphadev> Accessed: 2023-09-20



Conclusion



faster synthesis



faster sorting kernels



minimality proof

↑ Project on GitHub

